

# Software Development for A.I.

Title: **Health Insurance Premium Prediction using AWS Cloud**

### Important Libraries

1. Numpy
- NumPy is a library for the Python programming language, adding support for large, multi-dimensional arrays and matrices, along with a large collection of high-level mathematical functions to operate on these arrays.
2. Pandas
- Pandas is a software library written for the Python programming language for data manipulation and analysis. In particular, it offers data structures and operations for manipulating numerical tables and time series
3. Seaborn
- Seaborn is a Python data visualization library based on matplotlib. It provides a high-level interface for drawing attractive and informative statistical graphics.
4. scikit-learn
- Scikit-learn is an open source data analysis library, and the gold standard for Machine Learning (ML) in the Python ecosystem
5. TensorFlow
- TensorFlow is a free and open-source software library for machine learning and artificial intelligence. It can be used across a range of tasks but has a particular focus on training and inference of deep neural networks.

```
In [2]: #Importing the required libraries.
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [3]: # Reading the Dataset
insurance = pd.read_csv('insurance.csv')
```

```
In [4]: # Quick glance at the data
insurance
```

Out[4]:


	age	sex	bmi	children	smoker	region	expenses
0	19	female	27.9	0	yes	southwest	16884.92
1	18	male	33.8	1	no	southeast	1725.55
2	28	male	33.0	3	no	southeast	4449.46
3	33	male	22.7	0	no	northwest	21984.47
4	32	male	28.9	0	no	northwest	3866.86
...	...	...	...	...	...	...	...
1333	50	male	31.0	3	no	northwest	10600.55
1334	18	female	31.9	0	no	northeast	2205.98
1335	18	female	36.9	0	no	southeast	1629.83
1336	21	female	25.8	0	no	southwest	2007.95
1337	61	female	29.1	0	yes	northwest	29141.36

1338 rows × 7 columns

## Performing Exploratory Data Analysis

```
In [5]: #Checking for null values in the Dataset
sns.heatmap(insurance.isnull(), yticklabels = False, cbar = False, cmap="Blues")
```

Out[5]: <AxesSubplot: >



```
In [6]: # Grouping by region to see any relationship between region and charges
region = insurance.groupby(by='region').mean()
region
# It reveals that the south east region has higher BMI values when compared to northeast

/tmp/ipykernel_105/2587700523.py:2: FutureWarning: The default value of numeric_only in DataFrameGroupBy.mean is deprecated. In a future version, numeric_only will default to False. Either specify numeric_only or select only columns which should be valid for the function.
region = insurance.groupby(by='region').mean()
```

Out[6]:

	age	bmi	children	expenses
region				
northeast	39.268519	29.176235	1.046296	13406.384691
northwest	39.196923	29.201846	1.147692	12417.575169
southeast	38.939560	33.359341	1.049451	14735.411538
southwest	39.455385	30.596615	1.141538	12346.937908

```
In [7]: # Grouping by age
age = insurance.groupby(by='age').mean()
age

/tmp/ipykernel_105/2356038905.py:2: FutureWarning: The default value of numeric_only in DataFrameGroupBy.mean is deprecated. In a future version, numeric_only will default to False. Either specify numeric_only or select only columns which should be valid for the function.
age = insurance.groupby(by='age').mean()
```

Out[7]:

	bmi	children	expenses
--	-----	----------	----------

age			
18	31.333333	0.449275	7086.217971
19	28.598529	0.426471	9747.909706
20	30.627586	0.862069	10159.697931
21	28.189286	0.785714	4730.464286
22	31.092857	0.714286	10012.932857
23	31.460714	1.000000	12419.820357
24	29.142857	0.464286	10648.015714
25	29.689286	1.285714	9838.365000
26	29.435714	1.071429	6133.825714
27	29.342857	0.964286	12184.701429
28	29.482143	1.285714	9069.187500
29	29.374074	1.259259	10430.159630
30	30.555556	1.555556	12719.111111
31	29.922222	1.407407	10196.980000
32	31.588462	1.269231	9220.299615
33	31.165385	1.538462	12351.532308
34	30.273077	1.153846	11613.528462
35	31.392000	1.680000	11307.183200
36	29.368000	1.240000	12204.477600
37	31.216000	1.520000	18019.911600
38	29.004000	1.480000	8102.732800
39	29.908000	2.200000	11778.243600
40	30.144444	1.592593	11772.251481
41	31.518519	1.407407	9653.745556
42	30.337037	1.000000	13061.038519
43	30.207407	1.629630	19267.279630
44	30.848148	1.222222	15859.397037
45	29.782759	1.482759	14830.199310
46	31.341379	1.620690	14342.591379
47	30.655172	1.379310	17653.999655
48	31.927586	1.310345	14632.500000
49	30.314286	1.500000	12696.006071
50	31.134483	1.310345	15663.003103
51	31.731034	1.103448	15682.255517
52	32.941379	1.482759	18256.270345
53	30.371429	1.250000	16020.930357
54	31.232143	1.428571	18758.546429
55	31.950000	0.961538	16164.545000
56	31.600000	0.769231	15025.517692
57	30.838462	0.615385	16447.186154
58	32.728000	0.240000	13878.928000
59	30.576000	1.200000	18895.869600
60	30.330435	0.347826	21979.419130
61	32.552174	0.739130	22024.457391
62	32.347826	0.565217	19163.856087
63	31.930435	0.565217	19884.999130
64	32.986364	0.772727	23275.530909

In [8]:

```
# Conversion of categorical value to numerical value to avoid training issues with the machine learning model
insurance['sex'] = insurance['sex'].apply(lambda x: 0 if x == 'female' else 1)
insurance['smoker'] = insurance['smoker'].apply(lambda x: 1 if x == 'yes' else 0)
```

In [9]:

```
# Check unique values in 'region' column
insurance['region'].unique()
```

Out[9]: array(['southwest', 'southeast', 'northwest', 'northeast'], dtype=object)

```
In [10]: region_dummies = pd.get_dummies(insurance['region'], drop_first = True)
```

```
In [11]: insurance = pd.concat([insurance, region_dummies], axis = 1)
```

```
In [12]: insurance.describe()
```

Out[12]:


	age	sex	bmi	children	smoker	expenses	northwest	southeast	southwest
count	1338.000000	1338.000000	1338.000000	1338.000000	1338.000000	1338.000000	1338.000000	1338.000000	1338.000000
mean	39.207025	0.505232	30.665471	1.094918	0.204783	13270.422414	0.242900	0.272048	0.242900
std	14.049960	0.500160	6.098382	1.205493	0.403694	12110.011240	0.428995	0.445181	0.428995
min	18.000000	0.000000	16.000000	0.000000	0.000000	1121.870000	0.000000	0.000000	0.000000
25%	27.000000	0.000000	26.300000	0.000000	0.000000	4740.287500	0.000000	0.000000	0.000000
50%	39.000000	1.000000	30.400000	1.000000	0.000000	9382.030000	0.000000	0.000000	0.000000
75%	51.000000	1.000000	34.700000	2.000000	0.000000	16639.915000	0.000000	1.000000	0.000000
max	64.000000	1.000000	53.100000	5.000000	1.000000	63770.430000	1.000000	1.000000	1.000000

## Plotting the Dataset graph using Seaborn


1. Graphs showing the overview of aspects like Age, Sex, BMI, Children, Smoker, Expenses

```
In [13]: insurance[['age', 'sex', 'bmi', 'children', 'smoker', 'expenses']].hist(bins = 30, figsize = (20,20), color = 'r')
```


```
Out[13]: array([[<AxesSubplot: title={'center': 'age'}>,  
  <AxesSubplot: title={'center': 'sex'}>],  
  [<AxesSubplot: title={'center': 'bmi'}>,  
  <AxesSubplot: title={'center': 'children'}>],  
  [<AxesSubplot: title={'center': 'smoker'}>,  
  <AxesSubplot: title={'center': 'expenses'}>]], dtype=object)
```



```
In [14]: # Plotting age and expenses
sns.regplot(x = 'age', y = 'expenses', data = insurance)
plt.show()
```



```
In [15]: # Plotting BMI and expenses
sns.regplot(x = 'bmi', y = 'expenses', data = insurance)
plt.show()
```



```
In [16]: # Finding the corelation among them
correlation = insurance.corr()
correlation
```

/tmp/ipykernel\_105/2549767591.py:2: FutureWarning: The default value of numeric\_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric\_only to silence this warning.


```
correlation = insurance.corr()
```

Out[16]:

	age	sex	bmi	children	smoker	expenses	northwest	southeast	southwest
age	1.000000	-0.020856	0.109341	0.042469	-0.025019	0.299008	-0.000407	-0.011642	0.010016
sex	-0.020856	1.000000	0.046380	0.017163	0.076185	0.057292	-0.011156	0.017117	-0.004184
bmi	0.109341	0.046380	1.000000	0.012645	0.003968	0.198576	-0.135992	0.270144	-0.006398
children	0.042469	0.017163	0.012645	1.000000	0.007673	0.067998	0.024806	-0.023066	0.021914
smoker	-0.025019	0.076185	0.003968	0.007673	1.000000	0.787251	-0.036945	0.068498	-0.036945
expenses	0.299008	0.057292	0.198576	0.067998	0.787251	1.000000	-0.039905	0.073982	-0.043210
northwest	-0.000407	-0.011156	-0.135992	0.024806	-0.036945	-0.039905	1.000000	-0.346265	-0.320829
southeast	-0.011642	0.017117	0.270144	-0.023066	0.068498	0.073982	-0.346265	1.000000	-0.346265
southwest	0.010016	-0.004184	-0.006398	0.021914	-0.036945	-0.043210	-0.320829	-0.346265	1.000000

```
In [ ]: plt.figure(figsize = (10,10))
sns.heatmap(correlation , annot = True)
```

Out[ ]: <AxesSubplot: >



```
In [ ]: #Dropping Region coloumn
insurance.drop(['region'], axis = 1, inplace = True)
```

```
In [ ]: X = insurance.drop(columns =['expenses'])
y = insurance['expenses']
```

```
In [ ]: # Only take the numerical variables and scale them
X
```

Out[ ]:

	age	sex	bmi	children	smoker	northwest	southeast	southwest
0	19	0	27.9	0	1	0	0	1
1	18	1	33.8	1	0	0	1	0
2	28	1	33.0	3	0	0	1	0
3	33	1	22.7	0	0	1	0	0
4	32	1	28.9	0	0	1	0	0
...	...	...	...	...	...	...	...	...
1333	50	1	31.0	3	0	1	0	0
1334	18	0	31.9	0	0	0	0	0
1335	18	0	36.9	0	0	0	1	0
1336	21	0	25.8	0	0	0	0	1
1337	61	0	29.1	0	1	1	0	0

1338 rows × 8 columns

```
In [ ]: y
```

```
Out[ ]: 0      16884.92
        1      1725.55
        2      4449.46
        3     21984.47
        4      3866.86
        ...
    1333     10600.55
    1334      2205.98
    1335      1629.83
    1336      2007.95
    1337     29141.36
Name: expenses, Length: 1338, dtype: float64
```

```
In [31]: print('Shape of X:')
X.shape
```

Shape of X:

Out[31]: (1338, 8)

```
In [32]: print('Shape of Y:')
y.shape
```

Shape of Y:

Out[32]: (1338,)

```
In [33]: # Converting into Numpy Array
X = np.array(X).astype('float32')
y = np.array(y).astype('float32')
```

```
In [34]: y = y.reshape(-1,1)
```

```
In [35]: y.shape
```

Out[35]: (1338, 1)

```
In [36]: X
```

```
Out[36]: array([[19. ,  0. , 27.9, ...,  0. ,  0. ,  1. ],
                [18. ,  1. , 33.8, ...,  0. ,  1. ,  0. ],
                [28. ,  1. , 33. , ...,  0. ,  1. ,  0. ],
                ...,
                [18. ,  0. , 36.9, ...,  0. ,  1. ,  0. ],
                [21. ,  0. , 25.8, ...,  0. ,  0. ,  1. ],
                [61. ,  0. , 29.1, ...,  1. ,  0. ,  0. ]], dtype=float32)
```

# Model Building

```
In [37]: from sklearn.preprocessing import StandardScaler, MinMaxScaler

scaler_x = StandardScaler()
X = scaler_x.fit_transform(X)

scaler_y = StandardScaler()
y = scaler_y.fit_transform(y)
```

```
In [38]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3)
```

```
In [39]: X_train.shape
```

Out[39]: (936, 8)

```
In [40]: X_test.shape
```

Out[40]: (402, 8)

```
In [41]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, accuracy_score

regresssion_model_sklearn = LinearRegression()
regresssion_model_sklearn.fit(X_train, y_train)
```

Out[41]:

▼ LinearRegression

LinearRegression()

```
In [42]: regresssion_model_sklearn_accuracy = regresssion_model_sklearn.score(X_test, y_test)
print("regression model accuracy is " + str(regresssion_model_sklearn_accuracy*100))
```

regression model accuracy is 72.29441537204815

```
In [43]: y_predict = regresssion_model_sklearn.predict(X_test)
```

```
In [44]: y_predict_orig = scaler_y.inverse_transform(y_predict)
        y_test_orig = scaler_y.inverse_transform(y_test)

In [45]: k = X_test.shape[1]
        n = len(X_test)
        n

Out[45]: 402
```

## Calculating metrics for our regression model

```
In [46]: from sklearn.metrics import r2_score, mean_squared_error, mean_absolute_error
        from math import sqrt

        RMSE = float(format(np.sqrt(mean_squared_error(y_test_orig, y_predict_orig)),'.3f'))
        MSE = mean_squared_error(y_test_orig, y_predict_orig)

In [47]: MAE = mean_absolute_error(y_test_orig, y_predict_orig)
        r2 = r2_score(y_test_orig, y_predict_orig)
        adj_r2 = 1-(1-r2) * (n-1)/(n-k-1)

In [48]: print('RMSE =',RMSE, '\nMSE =',MSE, '\nMAE =',MAE, '\nR2 =', r2, '\nAdjusted R2 =', adj_r2)

RMSE = 6433.626
MSE = 41391544.0
MAE = 4358.8115
R2 = 0.7229441512449756
Adjusted R2 = 0.7173043375298606
```

## Using Artificial Neural Network

```
In [106... import tensorflow as tf
        from tensorflow import keras
        from tensorflow.keras.layers import Dense, Activation, Dropout, BatchNormalization
        from tensorflow.keras.optimizers import Adam

In [107... ANN_model = keras.Sequential()
        ANN_model.add(Dense(50, input_dim = 8))
        ANN_model.add(Activation('relu'))
        ANN_model.add(Dense(150))
        ANN_model.add(Activation('relu'))
        ANN_model.add(Dropout(0.5))
        ANN_model.add(Dense(150))
        ANN_model.add(Activation('relu'))
        ANN_model.add(Dropout(0.5))
        ANN_model.add(Dense(50))
        ANN_model.add(Activation('linear'))
        ANN_model.add(Dense(1))
        ANN_model.compile(loss = 'mse', optimizer = 'adam')
        ANN_model.summary()

Model: "sequential_2"

Layer (type)                Output Shape                Param #
=====
dense_10 (Dense)             (None, 50)                  450
activation_8 (Activation)    (None, 50)                   0
dense_11 (Dense)             (None, 150)                 7650
activation_9 (Activation)    (None, 150)                   0
dropout_4 (Dropout)          (None, 150)                   0
dense_12 (Dense)             (None, 150)                22650
activation_10 (Activation)   (None, 150)                   0
dropout_5 (Dropout)          (None, 150)                   0
dense_13 (Dense)             (None, 50)                  7550
activation_11 (Activation)   (None, 50)                   0
dense_14 (Dense)             (None, 1)                    51
=====
Total params: 38,351
Trainable params: 38,351
Non-trainable params: 0

In [108... ANN_model.compile(optimizer='Adam', loss='mean_squared_error')
```



```
epochs_hist = ANN_model.fit(X_train, y_train, epochs = 100, batch_size = 20, validation_split = 0.2)
```

Epoch 1/100  
38/38 [=====] - 3s 13ms/step - loss: 0.7304 - val\_loss: 0.4560  
Epoch 2/100  
38/38 [=====] - 0s 8ms/step - loss: 0.4466 - val\_loss: 0.3254  
Epoch 3/100  
38/38 [=====] - 0s 7ms/step - loss: 0.3499 - val\_loss: 0.2671  
Epoch 4/100  
38/38 [=====] - 0s 8ms/step - loss: 0.2580 - val\_loss: 0.2449  
Epoch 5/100  
38/38 [=====] - 0s 6ms/step - loss: 0.2391 - val\_loss: 0.2801  
Epoch 6/100  
38/38 [=====] - 0s 9ms/step - loss: 0.2409 - val\_loss: 0.2369  
Epoch 7/100  
38/38 [=====] - 0s 7ms/step - loss: 0.2400 - val\_loss: 0.2505  
Epoch 8/100  
38/38 [=====] - 0s 8ms/step - loss: 0.2139 - val\_loss: 0.2317  
Epoch 9/100  
38/38 [=====] - 0s 7ms/step - loss: 0.1907 - val\_loss: 0.2389  
Epoch 10/100  
38/38 [=====] - 0s 7ms/step - loss: 0.2154 - val\_loss: 0.2594  
Epoch 11/100  
38/38 [=====] - 0s 5ms/step - loss: 0.2221 - val\_loss: 0.2371  
Epoch 12/100  
38/38 [=====] - 0s 4ms/step - loss: 0.2163 - val\_loss: 0.2262  
Epoch 13/100  
38/38 [=====] - 0s 5ms/step - loss: 0.1900 - val\_loss: 0.2315  
Epoch 14/100  
38/38 [=====] - 0s 5ms/step - loss: 0.2024 - val\_loss: 0.2383  
Epoch 15/100  
38/38 [=====] - 0s 5ms/step - loss: 0.1876 - val\_loss: 0.2616  
Epoch 16/100  
38/38 [=====] - 0s 5ms/step - loss: 0.1945 - val\_loss: 0.2599  
Epoch 17/100  
38/38 [=====] - 0s 6ms/step - loss: 0.2015 - val\_loss: 0.2280  
Epoch 18/100  
38/38 [=====] - 0s 6ms/step - loss: 0.1746 - val\_loss: 0.2302  
Epoch 19/100  
38/38 [=====] - 0s 7ms/step - loss: 0.1910 - val\_loss: 0.2660  
Epoch 20/100  
38/38 [=====] - 0s 4ms/step - loss: 0.1824 - val\_loss: 0.2292  
Epoch 21/100  
38/38 [=====] - 0s 6ms/step - loss: 0.1899 - val\_loss: 0.2697  
Epoch 22/100  
38/38 [=====] - 0s 4ms/step - loss: 0.1831 - val\_loss: 0.2364  
Epoch 23/100  
38/38 [=====] - 0s 4ms/step - loss: 0.1776 - val\_loss: 0.2256  
Epoch 24/100  
38/38 [=====] - 0s 4ms/step - loss: 0.1792 - val\_loss: 0.2385  
Epoch 25/100  
38/38 [=====] - 0s 5ms/step - loss: 0.1743 - val\_loss: 0.2226  
Epoch 26/100  
38/38 [=====] - 0s 8ms/step - loss: 0.1773 - val\_loss: 0.2464  
Epoch 27/100  
38/38 [=====] - 0s 7ms/step - loss: 0.1771 - val\_loss: 0.2371  
Epoch 28/100  
38/38 [=====] - 0s 5ms/step - loss: 0.1713 - val\_loss: 0.2277  
Epoch 29/100  
38/38 [=====] - 0s 6ms/step - loss: 0.1578 - val\_loss: 0.2332  
Epoch 30/100  
38/38 [=====] - 0s 11ms/step - loss: 0.1628 - val\_loss: 0.2376  
Epoch 31/100  
38/38 [=====] - 0s 5ms/step - loss: 0.1727 - val\_loss: 0.2319  
Epoch 32/100  
38/38 [=====] - 0s 6ms/step - loss: 0.1724 - val\_loss: 0.2431  
Epoch 33/100  
38/38 [=====] - 0s 7ms/step - loss: 0.1824 - val\_loss: 0.2573  
Epoch 34/100  
38/38 [=====] - 0s 5ms/step - loss: 0.1725 - val\_loss: 0.2634  
Epoch 35/100  
38/38 [=====] - 0s 5ms/step - loss: 0.1680 - val\_loss: 0.2571  
Epoch 36/100  
38/38 [=====] - 0s 4ms/step - loss: 0.1679 - val\_loss: 0.2473  
Epoch 37/100  
38/38 [=====] - 0s 6ms/step - loss: 0.1589 - val\_loss: 0.2641  
Epoch 38/100  
38/38 [=====] - 0s 4ms/step - loss: 0.1700 - val\_loss: 0.2468  
Epoch 39/100  
38/38 [=====] - 0s 4ms/step - loss: 0.1626 - val\_loss: 0.2688  
Epoch 40/100  
38/38 [=====] - 0s 7ms/step - loss: 0.1545 - val\_loss: 0.2405  
Epoch 41/100  
38/38 [=====] - 0s 5ms/step - loss: 0.1672 - val\_loss: 0.2483  
Epoch 42/100  
38/38 [=====] - 0s 4ms/step - loss: 0.1549 - val\_loss: 0.2508  
Epoch 43/100  
38/38 [=====] - 0s 6ms/step - loss: 0.1476 - val\_loss: 0.2671  
Epoch 44/100  
38/38 [=====] - 0s 4ms/step - loss: 0.1509 - val\_loss: 0.2441  
Epoch 45/100  
38/38 [=====] - 0s 6ms/step - loss: 0.1449 - val\_loss: 0.2639

Epoch 46/100  
38/38 [=====] - 0s 4ms/step - loss: 0.1507 - val\_loss: 0.2396  
Epoch 47/100  
38/38 [=====] - 0s 4ms/step - loss: 0.1469 - val\_loss: 0.2594  
Epoch 48/100  
38/38 [=====] - 0s 5ms/step - loss: 0.1468 - val\_loss: 0.2399  
Epoch 49/100  
38/38 [=====] - 0s 5ms/step - loss: 0.1468 - val\_loss: 0.2372  
Epoch 50/100  
38/38 [=====] - 0s 6ms/step - loss: 0.1490 - val\_loss: 0.2361  
Epoch 51/100  
38/38 [=====] - 0s 5ms/step - loss: 0.1471 - val\_loss: 0.2685  
Epoch 52/100  
38/38 [=====] - 0s 6ms/step - loss: 0.1544 - val\_loss: 0.2461  
Epoch 53/100  
38/38 [=====] - 0s 5ms/step - loss: 0.1378 - val\_loss: 0.2461  
Epoch 54/100  
38/38 [=====] - 0s 6ms/step - loss: 0.1402 - val\_loss: 0.2313  
Epoch 55/100  
38/38 [=====] - 0s 7ms/step - loss: 0.1384 - val\_loss: 0.2526  
Epoch 56/100  
38/38 [=====] - 0s 6ms/step - loss: 0.1518 - val\_loss: 0.2530  
Epoch 57/100  
38/38 [=====] - 0s 4ms/step - loss: 0.1467 - val\_loss: 0.2474  
Epoch 58/100  
38/38 [=====] - 0s 4ms/step - loss: 0.1491 - val\_loss: 0.2343  
Epoch 59/100  
38/38 [=====] - 0s 4ms/step - loss: 0.1411 - val\_loss: 0.2373  
Epoch 60/100  
38/38 [=====] - 0s 5ms/step - loss: 0.1254 - val\_loss: 0.2591  
Epoch 61/100  
38/38 [=====] - 0s 4ms/step - loss: 0.1502 - val\_loss: 0.2416  
Epoch 62/100  
38/38 [=====] - 0s 5ms/step - loss: 0.1383 - val\_loss: 0.2400  
Epoch 63/100  
38/38 [=====] - 0s 4ms/step - loss: 0.1353 - val\_loss: 0.2368  
Epoch 64/100  
38/38 [=====] - 0s 4ms/step - loss: 0.1461 - val\_loss: 0.2644  
Epoch 65/100  
38/38 [=====] - 0s 4ms/step - loss: 0.1375 - val\_loss: 0.2285  
Epoch 66/100  
38/38 [=====] - 0s 5ms/step - loss: 0.1435 - val\_loss: 0.2368  
Epoch 67/100  
38/38 [=====] - 0s 5ms/step - loss: 0.1362 - val\_loss: 0.2404  
Epoch 68/100  
38/38 [=====] - 0s 5ms/step - loss: 0.1384 - val\_loss: 0.2491  
Epoch 69/100  
38/38 [=====] - 0s 5ms/step - loss: 0.1392 - val\_loss: 0.2535  
Epoch 70/100  
38/38 [=====] - 0s 5ms/step - loss: 0.1408 - val\_loss: 0.2424  
Epoch 71/100  
38/38 [=====] - 0s 6ms/step - loss: 0.1351 - val\_loss: 0.2411  
Epoch 72/100  
38/38 [=====] - 0s 7ms/step - loss: 0.1253 - val\_loss: 0.2396  
Epoch 73/100  
38/38 [=====] - 0s 6ms/step - loss: 0.1242 - val\_loss: 0.2428  
Epoch 74/100  
38/38 [=====] - 0s 7ms/step - loss: 0.1472 - val\_loss: 0.2834  
Epoch 75/100  
38/38 [=====] - 0s 5ms/step - loss: 0.1459 - val\_loss: 0.2404  
Epoch 76/100  
38/38 [=====] - 0s 7ms/step - loss: 0.1260 - val\_loss: 0.2402  
Epoch 77/100  
38/38 [=====] - 0s 8ms/step - loss: 0.1369 - val\_loss: 0.2677  
Epoch 78/100  
38/38 [=====] - 0s 6ms/step - loss: 0.1337 - val\_loss: 0.2439  
Epoch 79/100  
38/38 [=====] - 0s 7ms/step - loss: 0.1292 - val\_loss: 0.2493  
Epoch 80/100  
38/38 [=====] - 0s 6ms/step - loss: 0.1306 - val\_loss: 0.2390  
Epoch 81/100  
38/38 [=====] - 0s 6ms/step - loss: 0.1279 - val\_loss: 0.2368  
Epoch 82/100  
38/38 [=====] - 0s 8ms/step - loss: 0.1464 - val\_loss: 0.2760  
Epoch 83/100  
38/38 [=====] - 0s 5ms/step - loss: 0.1301 - val\_loss: 0.2448  
Epoch 84/100  
38/38 [=====] - 0s 6ms/step - loss: 0.1270 - val\_loss: 0.2354  
Epoch 85/100  
38/38 [=====] - 0s 5ms/step - loss: 0.1227 - val\_loss: 0.2530  
Epoch 86/100  
38/38 [=====] - 0s 5ms/step - loss: 0.1382 - val\_loss: 0.2448  
Epoch 87/100  
38/38 [=====] - 0s 5ms/step - loss: 0.1313 - val\_loss: 0.2354  
Epoch 88/100  
38/38 [=====] - 0s 5ms/step - loss: 0.1270 - val\_loss: 0.2410  
Epoch 89/100  
38/38 [=====] - 0s 5ms/step - loss: 0.1268 - val\_loss: 0.2453  
Epoch 90/100  
38/38 [=====] - 0s 5ms/step - loss: 0.1354 - val\_loss: 0.2420


Epoch 91/100  
38/38 [=====] - 0s 4ms/step - loss: 0.1393 - val\_loss: 0.2456  
Epoch 92/100  
38/38 [=====] - 0s 4ms/step - loss: 0.1268 - val\_loss: 0.2362  
Epoch 93/100  
38/38 [=====] - 0s 5ms/step - loss: 0.1223 - val\_loss: 0.2444  
Epoch 94/100  
38/38 [=====] - 0s 5ms/step - loss: 0.1316 - val\_loss: 0.2479  
Epoch 95/100  
38/38 [=====] - 0s 6ms/step - loss: 0.1263 - val\_loss: 0.2512  
Epoch 96/100  
38/38 [=====] - 0s 7ms/step - loss: 0.1223 - val\_loss: 0.2510  
Epoch 97/100  
38/38 [=====] - 0s 8ms/step - loss: 0.1175 - val\_loss: 0.2511  
Epoch 98/100  
38/38 [=====] - 0s 6ms/step - loss: 0.1329 - val\_loss: 0.2507  
Epoch 99/100  
38/38 [=====] - 0s 5ms/step - loss: 0.1478 - val\_loss: 0.2442  
Epoch 100/100  
38/38 [=====] - 0s 8ms/step - loss: 0.1264 - val\_loss: 0.2466

```
In [109... result = ANN_model.evaluate(X_test, y_test)
accuracy_ANN = 1 - result
print("Accuracy : {}".format(accuracy_ANN))
```

13/13 [=====] - 0s 2ms/step - loss: 0.1594  
Accuracy : 0.840594083070755

```
In [110... plt.plot(epochs_hist.history['loss'])
plt.plot(epochs_hist.history['val_loss'])
plt.title('Model Loss Progress During Training')
plt.xlabel('Epoch')
plt.ylabel('Training and Validation Loss')
plt.legend(['Training Loss', 'Validation Loss'])
```


Out[110]: <matplotlib.legend.Legend at 0x7f1f54220f10>



```
In [111... y_predict = ANN_model.predict(X_test)
plt.plot(y_test, y_predict, "^", color = 'r')
plt.xlabel('Model Predictions')
plt.ylabel('True Values')
```

13/13 [=====] - 0s 2ms/step

Out[111]: Text(0, 0.5, 'True Values')



## Deploying to Amazon Sagemaker

```
In [ ]: !pip install "sagemaker"
        !pip install boto3
```

```
In [ ]: import sagemaker
import boto3

sagemaker_session = sagemaker.Session()
#Creating a sagemaker session and passing bucket and prefix values

bucket = 'sagemaker-health1'
prefix = 'linear_learner'

role = sagemaker.get_execution_role()
print(role)
```

```
In [ ]: import io
import numpy as np
import sagemaker.amazon.common as smac # sagemaker common library

buf = io.BytesIO()
smac.write_numpy_to_dense_tensor(buf, X_train, y_train)
buf.seek(0)
```

```
In [ ]: import os

key = 'linear-train-data'
boto3.resource('s3').Bucket(bucket).Object(os.path.join(prefix, 'train', key)).upload_fileobj(buf)
s3_train_data = 's3://{}/{}/train/{}'.format(bucket, prefix, key)
print('uploaded training data location: {}'.format(s3_train_data))
```

```
In [ ]: output_location = 's3://{}/{}/output'.format(bucket, prefix)
print('Training artifacts will be uploaded to: {}'.format(output_location))
```

```
In [ ]: # This part of code will allow us to use training containers of sagemaker trained model.

from sagemaker.amazon.amazon_estimator import image_uris
container = image_uris.retrieve(region=boto3.Session().region_name, framework='linear-learner')
```

```
In [ ]: # Parameters are passed using Estimator function that is available to us.

linear = sagemaker.estimator.Estimator(container,
                                       role,
                                       instance_count = 1,
                                       instance_type = 'ml.c4.xlarge',
                                       output_path = output_location,
                                       sagemaker_session = sagemaker_session)

linear.set_hyperparameters(feature_dim = 8,
                           predictor_type = 'regressor',
                           mini_batch_size = 100,
                           epochs = 100,
                           num_models = 32,
                           loss = 'absolute_loss')

linear.fit({'train': s3_train_data})
```